

Lab Report 1

April 30th, 2017

Bradley Evans and Dharti Tarapara

CS153 Operating Systems

Part 1: Adding System Calls

Hello World

Implementing a hello world function and executing it in the xv6 shell.

The Hello World function is implemented by changing the following files.

- `defs.h`
 - A function prototype is added here.
 - `void hello(void);`
- `syscall.h`
 - Define a system call here.
 - `#define SYS_hello 22`
- `user.h`
 - Define a function prototype.
 - `int hello(void);`
- `proc.c`
 - Define the actual hello function.
 - `void hello(void) { cprintf("hello!\n"); }`
- `sysproc.c`
 - Define the system call here. This will simply call our `hello()` program in `proc.c`.
 - `int sys_hello(void) { hello(); return 0; }`
- `hello.c`
 - This is added to the root directory.

```
#include "types.h"
#include "stat.h"
#include "user.h"

int main(int argc, char * argv[]) {
    hello();
    exit(-1);

    return 0;
}
```

Implemented in this way, the user can now execute a `hello` from the command line. This will trigger a system call, which will eventually call `hello()` from `proc.c`, displaying the hello world message.

Editing `wait()` and `exit()`

The `exit()` and `wait()` functions were modified to take integers.

Each instance of `exit()` in user programs was changed to `exit(0)` to reflect the change in type of exit from `void` to `int`. `exit` (and, in the same way, `wait`) now return a status.

To make use of this new parameter, `exit(int)` was changed to the following:

```
void
exit(int status)
{
    struct proc *p;
    int fd;

    // cprintf("exit status %d", proc->status);

    if(proc == initproc)
        panic("init exiting");

    // Close all open files.
    for(fd = 0; fd < NOFILE; fd++){
        if(proc->ofile[fd]){
            fclose(proc->ofile[fd]);
            proc->ofile[fd] = 0;
        }
    }

    begin_op();
    input(proc->cwd);
    end_op();
    proc->cwd = 0;

    acquire(&ptable.lock);

    // Parent might be sleeping in wait().
    wakeup1(proc->parent);

    // Pass abandoned children to init.
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->parent == proc){
            p->parent = initproc;
            if(p->state == ZOMBIE)
                wakeup1(initproc);
        }
    }
}
```

```

proc->status = status; // MOD - 4/18
// Jump into the scheduler, never to return.
proc->state = ZOMBIE;
sched();
panic("zombie exit");
}

```

Also, `wait()` was changed to the following:

```

int
wait(int * status)
{
    struct proc *p;
    int havekids, pid;

    acquire(&ptable.lock);
    for(;;){
        // Scan through table looking for exited children.
        havekids = 0;
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->parent != proc)
                continue;
            havekids = 1;
            if(p->state == ZOMBIE){
                // Found one.

                // MOD - 4/29
                if (p->status != 0) {
                    *status = p->status;
                } else *status = 0;

                pid = p->pid;
                kfree(p->kstack);
                p->kstack = 0;
                freevm(p->pgdir);
                p->pid = 0;
                p->parent = 0;
                p->name[0] = 0;
                p->killed = 0;
                p->state = UNUSED;
                release(&ptable.lock);
                return pid;
            }
        }

        // No point waiting if we don't have any children.
        if(!havekids || proc->killed){
            release(&ptable.lock);
            return -1;
        }

        // Wait for children to exit. (See wakeup1 call in proc_exit.)
        sleep(proc, &ptable.lock); //DOC: wait-sleep
    }
}

```

Part 2: Schedulers

Implementing Priority Scheduling

Prior to this lab, xv6 used a *round robin* style scheduler. We modified it to use a *priority* scheduler, changing the scheduler code to the following.

```

void
scheduler(void)
{
    struct proc *p;
    int priority = 0; // hold priority value

    for(;;) {
        // Enable interrupts on this processor.
        sti();

        /* *** BEGIN MOD: 4/30 PRISCHED
        for(priority = 0; priority < maxpriority; priority++) {
            acquire(&ptable.lock);
            // Loop over process table looking for process to run.
            for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
                if(p->state != RUNNABLE)
                    continue;
                // PRISCHED: Now check and see if that process matches our current
                // priority level.
                if (p->priority == priority) {
                    // If it does, get it running.
                    proc = p;
                    switchvm(p);
                    p->state = RUNNING;
                    swtch(&cpu->scheduler, proc->context);
                    switchkvm();
                }
                // Process is done running for now.
                // It should have changed its p->state before coming back.
                proc = 0;
            }
            release(&ptable.lock);
        }*/

        // Loop over process table looking for process to run.
    }
}

```

```

    acquire(&ptable.lock);
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
        if(p->state != RUNNABLE) {
            continue;
        }

        release(&ptable.lock);

        priority = getprocpriority();

        acquire(&ptable.lock);

        if (priority < p->priority) {
            p->priority = priority;
        }
        // Switch to chosen process. It is the process's job
        // to release ptable.lock and then reacquire it
        // before jumping back to us.
        proc = p;
        switchvm(p);
        p->state = RUNNING;
        swtch(&cpu->scheduler, p->context);
        switchkvm();

        // Process is done running for now.
        // It should have changed its p->state before coming back.
        proc = 0;
    }
    release(&ptable.lock);
}
}

```

In addition, a priority variable was added to to struct proc to give variables a means of storing a priority code.

The processes need a way to change priority. A mutator is required. In proc.c this is defined as setpriority:

```

int
setpriority(int num)
{
    if (num < 0) {
        num = 0;
    } else if (num > 63) {
        num = 63;
    }
    priority

    proc->priority = num;

    return 0;
}

```

This is tied to a setpriority system call, added into the system in a similar way to the earlier hello call.

Also, a getprocpriority function will return the current seniormost priority level.

```

int
getprocpriority(void)
{
    struct proc *p;
    int priority = 65;

    // go through ptable
    acquire(&ptable.lock);
    for (p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
        // look for runnable process
        if(p->state != RUNNABLE) {
            continue;
        }

        if(priority == 0) {
            priority = p->priority;
        } else {
            if (p->priority < priority) {
                priority = p->priority;
            }
        }
    }

    release(&ptable.lock);

    return priority;
}

```

Priority Donation and Inheritance

A problem with priority schedulers is deadlock (more often called "unbounded priority inversion"), where lower-priority processes do not yield resources to a higher-priority process. A way to mitigate this is via *priority donation and inheritance*. If a higher priority task (H) attempts to access resource in use by a lower priority task (L), the lower priority task "inherits" a higher priority so that it can complete execution of some critical, uninterruptable portion of its programming. This prevents L from being pre-empted by a medium priority task M, which might then block H from executing.

To make this happen, you could use the setpriority mutator whenever a task enters the queue below a lower priority task. setpriority would manually change the lower priority's priority value to be equal to that of the high-priority task. When the high-priority task is allowed to execute, it can use setpriority to return the low-priority task to its original priority level.